


# Class 07: Machine Learning 1

Destiny (A16340362)

#Clustering We will start today's lab with clustering methods, in particular so-called K-means. The main functions for this in R is `kmeans()`

Let's try it on some made up data where we know what the answer should be

```
x <- rnorm(10000, mean=3)
hist(x)
```




60 points

```
tmp <- c(rnorm(30, mean=3), rnorm(30, -3))
x <- cbind(x=tmp, y=rev(tmp))
head(x)
```

```
      x      y
[1,] 4.120075 -2.247609
[2,] 4.215088 -2.501098
[3,] 2.908802 -4.237551
[4,] 3.305526 -4.682018
[5,] 2.236210 -3.053769
[6,] 2.223211 -1.950865
```

We can pass this to the base R `plot()` function for a quick plot

```
plot(x)
```



```
k <- kmeans(x, centers=2, nstart= 20)
k
```

K-means clustering with 2 clusters of sizes 30, 30

	x	y
1	2.762319	-2.869839
2	-2.869839	2.762319

[1] 1 2 2 2 2 2 2 2 2  
[39] 2

```
[1] 57.3778 57.3778
(between_SS / total_SS = 89.2 %)
```


```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

k\$size

```
k$cluster
```


k\$centers

```
plot(x, col= k$cluster, pch=16)
```



Q5. Cluster the data again with `kmeans()` into four groups and plot the results

```
k4 <- kmeans(x, centers=4, nstart= 20)  
plot(x, col= k4$cluster, pch=16)
```



K-means is very popular mostly because its fast and relatively straightforward to run and understand. It has a big limitation in that you need to tell it how many groups (k, or centers) you want. #Hierarchical clustering

The main function in base R is called `hclust()`. You have to pass it in a “distance matrix” not just your input data. (won’t work with just the data itself, need to pass it into the distance matrix)

You can generate a distance matrix with the `dist()` function.

```
hc <- hclust(dist(x))
hc
```


Call:

```
hclust(d = dist(x))
```

```
Cluster method : complete
Distance       : euclidean
Number of objects: 60
```

```
plot(hc)
```


## Cluster Dendrogram



To find the clusters (cluster membership vector) from a `hclust()` result we can “cut” the tree at a certain height that we like.

```
plot(hc)
abline(h=8, col="red")
```

## Cluster Dendrogram



dist(x)  
hclust (\*, "complete")


```
groups <- cutree(hc, h=8)
```

```
table(groups)
```

```
groups
 1  2
30 30
```

Q6. Plot our hclust results

```
plot(x, col= groups, pch=16)
```



#Principal Component Analysis

##PCA of UK food data

Read data showing the consumption in grams (per person, per week) of 17 different types of food stuff measured and averaged in the four countries of the United Kingdom

Let's see how PCA can help us but first we can try conventional analysis

```
url <- "https://tinyurl.com/UK-foods"
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

```
4.12007538325253 4.21508848293159 2.90880238997983 3.30552623758085
-2.247609 -2.501098 -4.237551 -4.682018
2.23620979966985 2.22321067565731
-3.053769 -1.950865
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
dim(x)
```



NULL

## Preview the first 6 rows

```
head(x)
```

```
4.12007538325253 4.21508848293159 2.90880238997983 3.30552623758085
-2.247609        -2.501098        -4.237551        -4.682018
2.23620979966985 2.22321067565731
-3.053769        -1.950865
```

```
x <- read.csv(url, row.names=1)
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

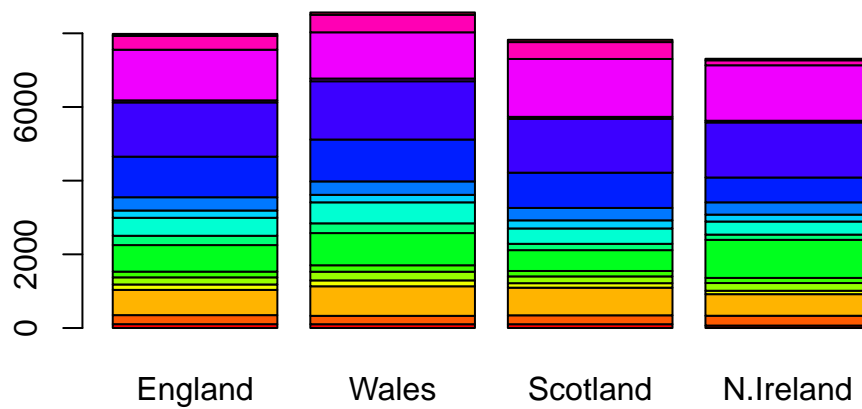
I prefer the `row.names=1` code. This approach is stronger than the `x[,-1]` approach because the code over writes the same object each time and it keeps getting rid of rows ever time that I run it.

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```




Q3: Changing what optional argument in the above `barplot()` function results in the following plot?

```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



```
pairs(x, col=rainbow(17), pch=16, cex=2)
```



Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

If a given point lies along the diagonal for a given plot means that the each country's consumption is about the same similarity, whereas if certain plots lie below the diagonal line this means dissimilarity and lower type of consumption.

Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

In this data set, the main differences between N. Ireland and the other countries in the UK is a bit more difficult to see. However, compared to the other countries in the UK and their graphs, each plot of N. Ireland vs the other UK countries has less of a straight diagonal line indicating that it's properly more dissimilar.

##Principal Component analysis (PCA)

PCA can help us make sense of these types of data sets. Let's see how it works.

The main function in "base" R is called `prcomp()`. In this case, we want to first take the transpose `t()` of our input `x` so the columns are the food types and the counties are the rows

```
head(t(x))
```

	Cheese	Carcass_meat	Other_meat	Fish	Fats_and_oils	Sugars
England	105	245	685	147	193	156
Wales	103	227	803	160	235	175
Scotland	103	242	750	122	184	147
N.Ireland	66	267	586	93	209	139
	Fresh_potatoes	Fresh_Veg	Other_Veg	Processed_potatoes		
England	720	253	488		198	
Wales	874	265	570		203	
Scotland	566	171	418		220	
N.Ireland	1033	143	355		187	
	Processed_Veg	Fresh_fruit	Cereals	Beverages	Soft_drinks	
England	360	1102	1472	57	1374	
Wales	365	1137	1582	73	1256	
Scotland	337	957	1462	53	1572	
N.Ireland	334	674	1494	47	1506	
	Alcoholic_drinks	Confectionery				
England	375	54				
Wales	475	64				

Scotland	458	62
N.Ireland	135	41


```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.


```
# Plot PC1 vs PC2
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x))
```



```
pca$x
```


	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-9.152022e-15
Wales	-240.52915	-224.646925	-56.475555	5.560040e-13
Scotland	-91.86934	286.081786	-44.415495	-6.638419e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.329771e-13

```
plot(pca$x[,1], pca$x[,2], col=c("orange", "red", "blue", "darkgreen"), pch=16)
```



Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x), col=c("orange", "red", "blue", "darkgreen"))
```




```
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
```

```
[1] 67 29 4 0
```

```
z <- summary(pca)
z$importance
```

	PC1	PC2	PC3	PC4
Standard deviation	324.15019	212.74780	73.87622	2.921348e-14
Proportion of Variance	0.67444	0.29052	0.03503	0.000000e+00
Cumulative Proportion	0.67444	0.96497	1.00000	1.000000e+00

```
barplot(v, xlab="Principal Component", ylab="Percent Variation")
```



The “loadings” tell us how much of the original variables (in our case the foods, contribute) to the new variables i.e. PCs


```
pca$rotation
```

	PC1	PC2	PC3	PC4
Cheese	-0.056955380	0.016012850	0.02394295	-0.409382587
Carcass_meat	0.047927628	0.013915823	0.06367111	0.729481922
Other_meat	-0.258916658	-0.015331138	-0.55384854	0.331001134
Fish	-0.084414983	-0.050754947	0.03906481	0.022375878
Fats_and_oils	-0.005193623	-0.095388656	-0.12522257	0.034512161
Sugars	-0.037620983	-0.043021699	-0.03605745	0.024943337
Fresh_potatoes	0.401402060	-0.715017078	-0.20668248	0.021396007
Fresh_Veg	-0.151849942	-0.144900268	0.21382237	0.001606882
Other_Veg	-0.243593729	-0.225450923	-0.05332841	0.031153231
Processed_potatoes	-0.026886233	0.042850761	-0.07364902	-0.017379680
Processed_Veg	-0.036488269	-0.045451802	0.05289191	0.021250980
Fresh_fruit	-0.632640898	-0.177740743	0.40012865	0.227657348
Cereals	-0.047702858	-0.212599678	-0.35884921	0.100043319
Beverages	-0.026187756	-0.030560542	-0.04135860	-0.018382072
Soft_drinks	0.232244140	0.555124311	-0.16942648	0.222319484
Alcoholic_drinks	-0.463968168	0.113536523	-0.49858320	-0.273126013



Confectionery                    -0.029650201   0.005949921   -0.05232164   0.001890737

```
## Lets focus on PC1 as it accounts for > 90% of variance
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,1], las=2 )
```



Q9: Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

The two food group features that are most prominent are fresh potatoes and soft drinks. PC1 captures the most variance through a certain point, but PC2 will capture the second most variance in comparison to PC1. In other words, PC2 will be able to capture the rest of the variance that PC1 wasn't able to capture.

```
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,2], las=2 )
```

